

KANTIPUR ENGINEERING COLLEGE

(Affiliated to Tribhuvan University)

Dhapakhel, Lalitpur



[Subject Code: CT707]

A MAJOR PROJECT MID-TERM REPORT ON FEATHERFIND : BIRD SPECIES IDENTIFICATION FROM AUDIO

Submitted by:

Gaurav Giri [Kan077bct034]

Iza K.C. [Kan077bct039]

Prajwal Khatiwada [Kan077bct056]

Samrat Kumar Adhikari [Kan077bct074]

A MAJOR PROJECT SUBMITTED IN PARTIAL FULFILLMENT OF THE REQUIREMENT FOR THE DEGREE OF BACHELOR IN COMPUTER ENGINEERING

Submitted to:

Department of Computer and Electronics Engineering

February, 2025

**FEATHERFIND : BIRD SPECIES IDENTIFICATION
FROM AUDIO**

Submitted by:

Gaurav Giri	[Kan077bct034]
Iza K.C.	[Kan077bct039]
Prajwal Khatiwada	[Kan077bct056]
Samrat Kumar Adhikari	[Kan077bct074]

**A MAJOR PROJECT SUBMITTED IN PARTIAL
FULFILLMENT OF THE REQUIREMENT FOR THE DEGREE
OF BACHELOR IN COMPUTER ENGINEERING**

Submitted to:

**Department of Computer and Electronics Engineering
Kantipur Engineering College
Dhapakhel, Lalitpur**

February, 2025

ABSTRACT

TABLE OF CONTENTS

Abstract	i
List Of Figures	vi
List Of Tables	vi
Abbreviations	vii
1 Introduction	1
1.1 Problem Statement	1
1.2 Objectives	2
1.3 Application Scope	2
1.4 Features	3
1.5 Feasibility Study	3
1.5.1 Economic Feasibility	4
1.5.2 Technical Feasibility	4
1.5.3 Operational Feasibility	4
1.5.4 Schedule Feasibility	5
1.6 System Requirements	6
1.6.1 Development Requirements	6
1.6.2 Deployment Requirements	6
2 Literature Review	7
2.1 Related Works	7
2.2 Related Research	8
3 Methodology	14
3.1 System Overview	14
3.2 Working Mechanism for Bird Sound Detection	15
3.2.1 Data Collection	16
3.2.2 Data Transformation	18
3.2.3 Feature Extraction	19
3.2.4 Model Architecture	20
3.2.5 Training Process	23
3.2.6 Evaluation	24
3.3 Working Mechanism for Bird Species Classification	26

3.3.1	Data Collection	26
3.3.2	Data Augmentation	27
3.3.3	Data Preprocessing	28
3.3.4	Feature Extraction	29
3.3.5	Model Architecture	29
3.3.6	Evaluation Metrics	36
3.4	Feature Extraction	37
3.4.1	Mel Spectrogram	37
3.4.2	Mel-Frequency Cepstral Coefficients (MFCC)	39
3.5	Hyperparameter Optimization using Genetic Algorithm	40
4	Results and Analysis	42
4.1	Bird Sound Detection	42
4.2	Bird Species Classification	45
4.2.1	Classification using CNN-LSTM	45
4.2.2	Classification using EfficientNetB3	48
4.3	Hyperparameter Optimization Using Genetic Algorithim	51
4.3.1	Hyperparameters for CNN-LSTM	51
4.3.2	Hyperparameters for Efficient-Net	55
4.4	Mobile Application	58
5	Discussion	64
5.1	Achievements	64
5.2	Limitations	65
5.3	Future Improvements	65
References		65

LIST OF FIGURES

1.1	Gantt Chart	5
3.1	System Overview	14
3.2	Block diagram for the working mechanism for Bird Detection	15
3.3	Dataset distribution for freefield1010	16
3.4	Dataset distribution for warblrb10k	17
3.5	Distribution of the merged dataset	17
3.6	Modified ResNet-50 architecture for Bird Sound Detection.	21
3.7	Block diagram for the working mechanism for Bird Species Identification	27
3.8	Dataset before performing augmentation	28
3.9	Dataset after performing augmentation	29
3.10	Modified EfficientNetB3 architecture for Bird Species Classification. . .	31
3.11	Mel Spectrogram image for a sample audio	38
3.12	Feature Extraction Using Spectrogram and MFCC	39
3.13	MFCC image for a sample audio	40
4.1	Accuracy curves	43
4.2	Loss curves	43
4.3	Training and validation curves for Bird Sound Detection Model	43
4.4	Confusion Matrix for Bird Sound Detection Model	44
4.5	ROC Curve of the detection model	44
4.6	CNN-LSTM Model Evaluation	45
4.7	CNN-LSTM Confusion Matrix	46
4.8	CNN-LSTM Classification Report	47
4.9	EfficientNetB3 Model Evaluation	48
4.10	EfficientNetB3 Confusion Matrix	49
4.11	EfficientNetB3 Classification Report	50
4.12	Accuracy and Loss for CNN-LSTM in generation 1	52
4.13	Accuracy and Loss for CNN-LSTM in generation 2	53
4.14	Accuracy and Loss for CNN-LSTM in generation 3	53
4.15	Accuracy and Loss for CNN-LSTM in generation 4	54
4.16	Accuracy and Loss for CNN-LSTM in generation 5	54

4.17	Accuracy and Loss for EfficientNetB3 in generation 1	56
4.18	Accuracy and Loss for EfficientNetB3 in generation 2	56
4.19	Accuracy and Loss for EfficientNetB3 in generation 3	57
4.20	Accuracy and Loss for EfficientNetB3 in generation 4	57
4.21	Accuracy and Loss for EfficientNetB3 in generation 5	57
4.22	Home Page of FeatherFind.	59
4.23	Mapped Birds in FeatherFind.	60
4.24	Recording using FeatherFind.	61
4.25	Confirmation for Audio Recognition.	62
4.26	Bird Species Identification using FeatherFind.	63

LIST OF TABLES

1.1	Development Requirements	6
1.2	Deployment Requirements	6
3.1	Training Hyperparameters	25
3.2	Initial Training Hyperparameters	34
3.3	Initial Training Hyperparameters for CNN+LSTM	36
4.1	Performance Metrics for Bird Sound Detection Model	42
4.2	Hyperparameter search space for CNN-LSTM	51
4.3	Hyperparameter configurations and test accuracies for CNN-LSTM . . .	54
4.4	Hyperparameter search space for the EfficientNetB3	55
4.5	Hyperparameter configurations for Efficient-NetB3	58

ABBREVIATIONS

CNN	Convolutional Neural Network
DCNN	Deep Convolutional Neural Network
DCT	Discrete Cosine Transform
DFT	Discrete Fourier Transform
GA	Genetic Algorithm
GPS	Global Positioning System
MAP	Mean Average Precision
MLP	Multilayer Perceptron
RGB	Reg Green Blue
ROC	Receiver Operating characteristics
RNN	Recurrent Neural Network
STFT	Short-Time Fourier Transform

CHAPTER 1

INTRODUCTION

Globally, the avian kingdom is vast, with over 11,000 species, showing nature's complexity and evolutionary skill. This number, from the International Ornithological Committee as of April 2023, shows the great bird diversity, with each species having its own role and story.[1]

In Nepal, a country known for its rich nature and varied ecosystems, from the lowland Terai to the high Himalayas, there are more than 887 bird species, according to the Himalayan Nature organization. This is over 8% of the world's known bird species, a big number given Nepal's small size.[2]

Many of these species are endangered due to habitat loss, climate change, and human activities. The National Red List of Nepal's Birds identifies 168 nationally threatened bird species, including 68 Critically Endangered, 38 Endangered, and 62 Vulnerable species, as detailed in a publication by the Journal of Threatened Taxa.[3]

The situation of these endangered species shows the need for conservation efforts. Technologies such as audio recognition provide new methods for identifying and monitoring bird populations. By analyzing bird sounds and pictures, researchers can better understand species distribution, behavior, and threats. These technologies not only help conserve endangered species but also support broader biodiversity research.

1.1 Problem Statement

In Nepal, a hotspot of avian biodiversity, accurately identifying and classifying bird species, particularly those that are endangered, is a critical yet complex task. Traditional observation methods are limited by the vast geographical and ecological diversity of the region, making it challenging to monitor and protect these birds effectively. The necessity for precise identification is paramount for conservation efforts aimed at maintaining ecosystem balance. To address this, there is a pressing need for a method that can overcome these constraints by leveraging advanced technologies capable of distin-

guishing between the myriad of bird calls and songs, as well as visual markers through image classification. Such a method promises to automate the identification process, enhancing accuracy and efficiency in monitoring endangered species.

1.2 Objectives

- i To develop and implement an integrated technological solution that utilizes advanced audio recognition technique for the accurate identification and monitoring of bird species in Nepal, with an emphasis on endangered species.

1.3 Application Scope

1. Conservation Efforts:

This system will enhance conservation by enabling accurate monitoring of bird populations, helping track endangered species and take a step towards habitat protection.

2. Biodiversity Monitoring:

Automated identification will aid biodiversity monitoring by processing large datasets, helping detect species distribution on bird communities.

3. Ecological Research:

Researchers can use the system to study bird migration, and habitat use, providing crucial data for modeling ecosystems and understanding ecological interactions.

4. Environmental Education and Awareness:

Integrated into educational programs, this tool will raise public awareness about biodiversity and conservation, engaging students and scientists in bird identification.

5. Bird viewing:

Bird enthusiasts will benefit from this system as it will enhance bird watching experiences by providing instant identification of bird species

1.4 Features

1. Species Identification Using Audio:

The app allows users to record bird sounds in real-time using their device's microphone or upload pre-recorded audio files. Advanced noise filtering techniques isolate bird calls from background noise, and sound wave analysis helps in identifying distinct frequency patterns. Machine learning algorithms, trained on a vast database of bird calls, match the recorded sound to identify the bird species accurately.

2. Mapping Identified Bird Habitat:

The app tags the location of identified birds using GPS, providing detailed habitat information typical of each species. Integrating with mapping services, it displays bird sightings on an interactive map, generating heat maps to show species density and distribution. Additionally, it tracks and visualizes bird migration patterns over time, helping users understand seasonal movements.

3. Provide Description About the Birds:

For each identified bird species, the app offers detailed profiles that include scientific and common names, physical descriptions, and conservation status. It also provides audio and visual media for reference, along with information on the bird's behavior, diet, and typical habitats, enriching the user's understanding of the species.

1.5 Feasibility Study

Before implementation of project design, the feasibility analysis of the project must be done to move any further. The feasibility analysis of the project gives an idea on how the project will perform and its impact in the real world scenario. So, it is of utmost importance.

1.5.1 Economic Feasibility

Our system is economically achievable as a result of the development of several tools, libraries, and frameworks. Since all the software required to construct it is free and readily available online, this project is incredibly cost-effective. Only time and effort are needed to create a worthwhile, genuinely passive system. The project doesn't come at a substantial cost. From an economic standpoint, the project appears successful in this sense.

1.5.2 Technical Feasibility

The software needed to implement a project can be downloaded from a wide variety of online resources. Technically speaking, the project is feasible as the necessary software is easily available. We were able to learn the information we required for the project through a variety of online sources, including classes. All the libraries and data are accessible online for free because this project does not require any licensing costs. It is technically possible if one has the necessary information and resources.

1.5.3 Operational Feasibility

The project aims to enhance bird species identification through audio classification, making bird sound recognition more accessible and efficient. This solution is particularly beneficial for ornithologists, bird watchers, and environmental researchers who require accurate and quick identification of bird species based on their calls. The project leverages advanced audio signal processing and deep learning algorithms to classify bird sounds, ensuring high accuracy and reliability. Given the widespread availability of mobile devices and recording equipment, the project is operationally feasible, as it can be easily integrated into existing workflows and tools used by bird enthusiasts and professionals. By providing an efficient method for bird sound classification, the project supports a sizable community interested in avian studies and conservation, ensuring practical applicability and ease of use.

1.5.4 Schedule Feasibility

The workload of the project is divided amongst the project members. The scheduling is done according to an incremental model where different modules are planned to be assigned to the group members. So, the project fulfills the schedule feasibility requirements.

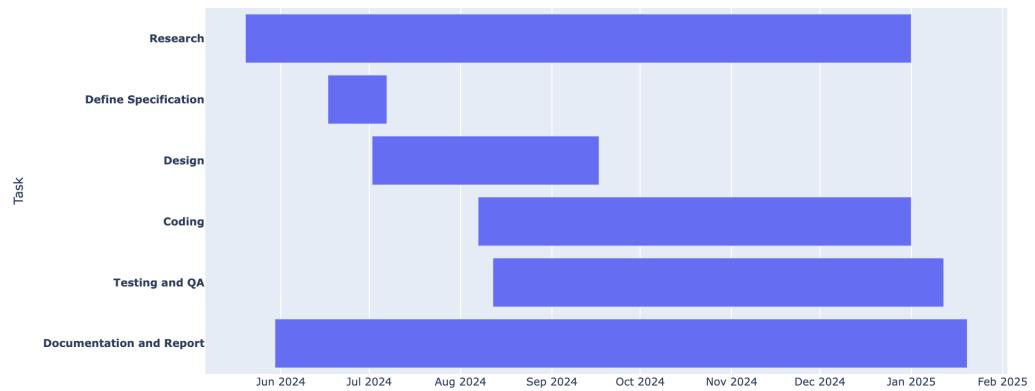


Figure 1.1: Gantt Chart

1.6 System Requirements

1.6.1 Development Requirements

Table 1.1: Development Requirements

Software Requirements	Hardware Requirements
Programming Language: Python, Dart, Javascript	RAM: \geq 8 GB Megapixels
Design Tools: Figma, Canva, Draw.io	CPU: i5 10th
Libraries: Librosa, PyAudio, Pytorch	GPU: P100 (Recommended)
Framework: Flutter, Django RestFramework	Storage: \geq 50 GB
IDEs: VSCode, Android Studio	

1.6.2 Deployment Requirements

Table 1.2: Deployment Requirements

Software Requirements	Hardware Requirements
Android: \geq 10	RAM: $>$ 4 GB
Read/Write FileSystem	Storage: \geq 20 GB
Internet Accessibility	Recording Quality \geq 256 Kbps, 48 KHz
Database: MySQL	

CHAPTER 2

LITERATURE REVIEW

This literature review explores the progression of methodologies and technologies in the field, with a particular focus on the use of Convolutional Neural Networks (CNN) and Long Short-Term Memory (LSTM) networks for audio-based bird species identification. The review also examines the challenges associated with dataset quality and diversity, and the innovative strategies employed to address these issues, providing a comprehensive overview of the current state of research and future directions in avian bioacoustics.

2.1 Related Works

BirdNET is a cutting-edge research platform developed through collaboration between the K. Lisa Yang Center for Conservation Bioacoustics at the Cornell Lab of Ornithology and the Chair of Media Informatics at Chemnitz University of Technology. Its primary aim is to detect and classify bird sounds using machine learning technologies, serving both experts and citizen scientists in their efforts to monitor and protect bird populations.

BirdNET can identify around 3,000 of the world's most common bird species, with plans to expand this number. Features such as a live submissions map and a Twitter bot are included to engage the community and share real-time data. The project is supported by donations and collaborations, offering opportunities for researchers and developers to contribute to its growth. BirdNET serves as an invaluable tool for bird enthusiasts, conservationists, and biologists alike, providing innovative solutions for large-scale acoustic monitoring and contributing to the conservation and understanding of avian biodiversity.

The BirdCLEF 2023 competition on Kaggle is a significant data science challenge that falls under the broader LifeCLEF initiative, aimed at pushing the boundaries of species identification and biodiversity monitoring through technological innovation. This par-

ticular competition focuses on the development of machine learning models that can identify bird species based on audio recordings. It presents a complex and realistic challenge due to the diversity of the audio recordings, which are collected from various environments and feature a wide range of bird species.

2.2 Related Research

The study[4] "Audio Classifier for Automatic Identification of Endangered Bird Species of Nepal" focuses on using deep learning techniques to identify endangered bird species from audio recordings. The dataset, collected from xeno-canto.org, comprises 2215 audio recordings of 41 bird species, 38 of which are endangered. This dataset was expanded to 6733 recordings through 10-second audio splitting and Gaussian noise augmentation, with 5407 recordings used for training, 639 for validation, and 687 for testing. The methodology involved handling imbalanced class distribution through data augmentation, employing Mel spectrograms and Mel-Frequency Cepstral Coefficients (MFCCs) for feature extraction, and developing a custom Convolutional Neural Network (CNN) model and an EfficientNet model. The hyperparameters of these models were optimized using a genetic algorithm. The Mel spectrograms were created using Short-Time Fourier Transform, converting amplitudes to decibel scale, and applying Mel filter banks to the spectrograms. Similarly, MFCCs were derived by framing the audio signals, applying Discrete Fourier Transform, logarithmic scaling, Mel scaling, and Discrete Cosine Transform. The EfficientNet architecture utilized compound scaling for network depth, width, and resolution. The findings indicated that the proposed approach achieved satisfactory results in classifying the bird species. Model I, using Mel Spectrogram and EfficientNet, achieved an F1-score of 79%, while Model II, using Mel Spectrogram and Custom CNN, achieved 64%, and Model III, using MFCC and EfficientNet, achieved 72%. However, limitations include the relatively small dataset size and the need for further enhancement in model robustness and accuracy.

Sevilla et al.[5] introduced an innovative approach to bird sound classification with their study 'Audio Bird Classification with Inception-v4 extended with Time and Time-Frequency Attention Mechanisms'. The datasets employed include various bird sounds,

prominently from the BirdClef2017 challenge, consisting of 1500 bird species recordings. The methodology revolves around treating bird sound classification as an image classification problem through transfer learning. The Inception-v4 model, initially pre-trained on ImageNet, was adapted to process time-frequency representations of bird sounds by converting these sounds into RGB images using three log-spectrograms generated via fast Fourier transform at different scales (128, 512, 2048 bins). The findings demonstrate that the model, termed ‘Soundception’, integrates time and time-frequency attention mechanisms effectively, significantly improving classification accuracy. The results highlight Soundception’s outstanding performance, achieving a mean average precision (MAP) of 0.714 in classifying 1500 bird species, 0.616 MAP for background species, and 0.288 MAP for soundscapes with time-codes, making it the top model in the BirdClef2017 challenge across multiple tasks. However, limitations include the incomplete convergence of the model due to computational constraints and the extensive GPU resources required for training. The paper concludes with a discussion on future improvements, such as exploring different scalable optimizations and incorporating stacked GRU layers for better audio-to-image representation learning, underscoring the potential of transfer learning from advanced image classification models to acoustic domains.

The study, conducted by Chandu B et al.[6], outlines a robust methodology for identifying bird species from audio recordings, leveraging a combination of meticulously curated datasets and machine learning techniques. The dataset was manually compiledThe working mechanism for bird identification from audio utilizes the dataset compiled from Xeno Canto, housing a diverse collection of avian vocalizations. Our methodology revolves around the utilization of EfficientNet, a state-of-the-art convolutional neural network architecture known for its balance between accuracy and efficiency.

The primary objective of this study is to achieve high levels of accuracy in identifying a broad spectrum of 41 distinct bird species. Here lies the detailed explanation of the methodology, from the working mechanism to model training. from both local recordings and online resources such as xeno-canto.org, which apart from bird songs also contains ambient noise and human voices to simulate real-world conditions. Pre-processing

techniques including pre-emphasis, framing, silence removal, and reconstruction were applied to the audio clips to enhance the relevant frequency components and eliminate unnecessary noise, ensuring the purity of the dataset. Spectrograms of these processed clips were generated and used as input for training a convolutional neural network (CNN), specifically AlexNet, chosen for its high accuracy in image classification tasks. Through transfer learning, AlexNet was adapted to recognize bird species from the spectrograms, achieving a classification accuracy of 97% in controlled environments. However, recognizing the variability of real-world conditions, the researchers retrained the model with datasets containing ambient noise, achieving a real-time classification accuracy of 91%. Despite these promising results, the study acknowledges limitations such as the relatively small size of the dataset and the need for further tuning of performance parameters to improve robustness.

In[7] ‘An Ensemble of Convolutional Neural Networks for Audio Classification’ delves into a comprehensive study on CNN classification using different architectures, data augmentation techniques, and audio signal representations, aimed at enhancing audio classification tasks across various datasets. The study employs three datasets: BIRDZ, CAT, and ESC-50, each offering unique challenges in audio classification. The methodology involves training five convolutional neural networks (CNNs) with four audio representations combined with six different data augmentation methods, resulting in thirty-five subtypes of ensembles. The audio representations include techniques such as the Discrete Gabor Transform (DGT), Waveform Similarity OverLap Add (WSOLA), and Phase Vocoder. The data augmentation methods encompass procedures like short spectrogram augmentation, random time shift, and frequency masking. The CNN architectures are pre-trained models fine-tuned with these augmented datasets to boost classification accuracy. The findings reveal that the ensemble method outperforms standalone networks, achieving 97% accuracy on the BIRDZ dataset, 90.51% on the CAT dataset, andThe working mechanism for bird identification from audio utilizes the dataset compiled from Xeno Canto, housing a diverse collection of avian vocalizations. Our methodology revolves around the utilization of EfficientNet, a state-of-the-art convolutional neural network architecture known for its balance between accuracy and ef-

ficiency.

The primary objective of this study is to achieve high levels of accuracy in identifying a broad spectrum of 41 distinct bird species. Here lies the detailed explanation of the methodology, from the working mechanism to model training. 88.65% on the ESC-50 dataset. The study also highlights that the best-performing CNNs are VGG16 and VGG19, with DGT as the most effective signal representation. However, the study acknowledges limitations, such as the computational cost of training ensembles and the variability in performance across different augmentation techniques.

In[8] ‘Analysis of bird call datasets sourced from Xeno-Canto’, comprising 72,172 samples from 264 bird species in 16-bit wav format with a 16 kHz sampling rate. The methodology involved preprocessing the audio data to filter out low-frequency noise and normalize signal amplitude, followed by generating Mel Spectrograms and Mel-Frequency Cepstral Coefficients (MFCCs) as inputs for deep learning models. The Mel Spectrograms were produced using discrete Fourier transform (DFT), and the MFCCs were derived by applying discrete cosine transform (DCT) to the Mel Spectrogram. The study employed various metrics to evaluate the performance of these methods, including ROC analysis to visualize model effectiveness. Findings indicated that the proposed models showed significant promise in identifying bird species from their calls, with improvements in classification accuracy compared to previous approaches. However, limitations were noted, including potential biases in the dataset due to uneven sample distribution across species and the challenge of background noise affecting signal quality. Future work suggested enhancing noise reduction techniques and exploring more sophisticated neural network architectures to further improve model robustness and accuracy.

The study[9] by Pahuja et al. conducted an in-depth analysis of bird species recognition through acoustic monitoring, utilizing a robust dataset of bird sound samples, meticulously annotated and validated for accuracy. The dataset, referred to as SD, comprises multispecies bird sound recordings, each labeled with species name and sample ID,

along with corresponding metadata, providing a comprehensive foundation for model training and evaluation. Methodologically, the research employed a spectrogram-based feature extraction approach, leveraging Short-Time Fourier Transform (STFT) to capture the intricate temporal and spectral characteristics of bird sounds. This was followed by the application of a Multilayer Perceptron (MLP) classifier to distinguish between different bird species. The findings reveal that the proposed model achieved high recognition accuracy, with some species being identified with perfect precision, recall, and accuracy (100%), though the performance varied across species, with a few showing lower recognition rates (86.9%) and precision/recall values ranging between 50-75%. The results demonstrated an overall classification accuracy of 96%, with cross-validation accuracy standing at 81.4%, highlighting the model's robustness yet indicating room for improvement in generalizability across diverse datasets. Despite the promising results, the study acknowledges several limitations, including the variability in recognition accuracy among different species and the potential influence of environmental noise on model performance. Future work is suggested to explore feature and model fusion techniques, integrate the model with cloud-based systems for real-time recognition, and expand the dataset to include a broader range of bird species to enhance the model's applicability and accuracy in practical scenarios.

The dataset used in this study[10] comprises recordings labeled by species from California and Nevada, USA. It includes 91 species, with 30 audio samples per species, amounting to a total of 2,730 MP3 files. The methodology adopted involves three main steps: pre-processing, feature extraction, and deep learning modeling. For feature extraction, MFCCs were obtained using the Python library `python_speech_features`, with parameters such as sample rate, 13 cepstrum coefficients, 26 filterbank filters, and an FFT size of 512. Mel spectrograms were extracted using the Librosa library, employing parameters such as a sample rate, an FFT window size of 2048, a hop length of 512, and 128 Mel bands. In the deep learning modeling, CNNs and LSTMs were compared for their effectiveness in classifying bird sounds. CNNs demonstrated superior training accuracy with Mel spectrogram features, achieving 99.05% and 98.76% accuracy for 3-second and 1.5-second spectrograms. The working mechanism for bird identification from

audio utilizes the dataset compiled from Xeno Canto, housing a diverse collection of avian vocalizations. Our methodology revolves around the utilization of EfficientNet, a state-of-the-art convolutional neural network architecture known for its balance between accuracy and efficiency.

The primary objective of this study is to achieve high levels of accuracy in identifying a broad spectrum of 41 distinct bird species. Here lies the detailed explanation of the methodology, from the working mechanism to model training. In contrast, LSTMs achieved lower training accuracies of 75.85% and 73.29% under similar conditions. These results highlight the superior ability of CNNs to leverage the spatial and frequency-related patterns in Mel spectrograms for accurate bird species classification.

Finally, in the study[11], 'Acoustic Bird Detection with Deep Convolution Neural Network', DCNNs pretrained on ImageNet, have been used for bioacoustic classification, with preprocessing and augmentation playing crucial roles in model performance. The preprocessing pipeline typically involves applying a shallow high-pass filter ($Q = 0.707$) with a 2 kHz cutoff, resampling to 22,050 Hz, and extracting around 4-second audio chunks, which are transformed into mel spectrograms with 310 mel bands (160-10,300 Hz). Further processing includes removing extreme frequencies, normalizing power spectrograms to decibel units, resizing to 224*224, and converting grayscale spectrograms to RGB for compatibility with ResNet. To improve generalization, data augmentation techniques such as jittering chunk duration, extracting chunks from random positions, adding noise from unrelated audio files, and applying random amplitude scaling are employed. Additional augmentations in the frequency domain include frequency shifting/stretching, piecewise time/frequency resizing, and color jittering (brightness, contrast, saturation, hue). Among these, the most effective methods are noise addition from random files, piecewise time and frequency stretching, and time interval dropout, which enhance robustness by simulating real-world variations in bird vocalizations.

CHAPTER 3

METHODOLOGY

This chapter describes the overall system including the bird sound detection and bird species classification.

3.1 System Overview




Figure 3.1: System Overview

Our project leverages a mobile application to capture audio recordings, which are then processed to identify bird species. The mobile app records the sound and visualizes it with waveforms during the recording process. Once the audio is captured, it is sent to our backend system for further analysis.

This model determines whether the recorded audio contains bird sounds. If bird sounds are detected, the audio is then forwarded to the bird species classifier model. The classifier identifies the specific bird species present in the recording. If the bird sound detection model does not detect any bird sounds, the recorded audio is labeled accordingly, indicating the absence of bird sounds.

Upon successful classification, the system maps the identified bird species to the location where the audio was recorded, utilizing the device's GPS data. This approach

ensures that only relevant audio recordings are processed for species classification, enhancing the accuracy and efficiency of the system. The integration of Django, Django Rest Framework, and MySQL provides a robust and scalable backend infrastructure to support the application's functionality as shown in Figure 3.1.

3.2 Working Mechanism for Bird Sound Detection

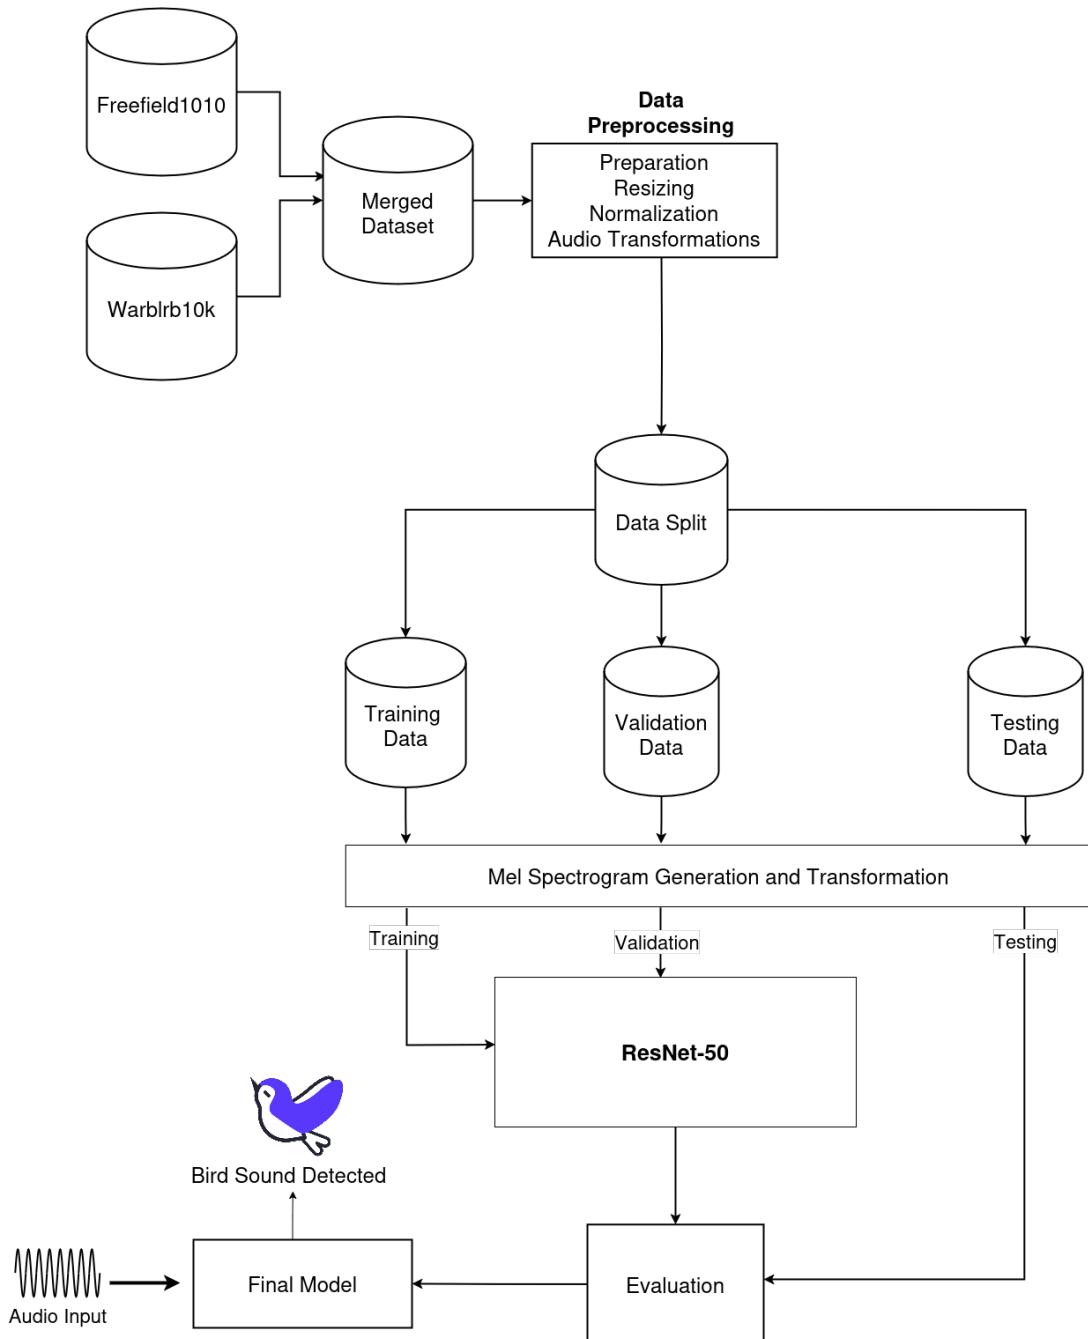


Figure 3.2: Block diagram for the working mechanism for Bird Detection

3.2.1 Data Collection

The dataset used for the Bird Sound Detection Model includes Field recordings, worldwide ("freefield1010") - a collection of 7,690 excerpts from field recordings around the world, gathered by the FreeSound project, and then standardised for research. This collection is very diverse in location and environment, and for the BAD Challenge, it has been annotated for the presence/absence of birds.

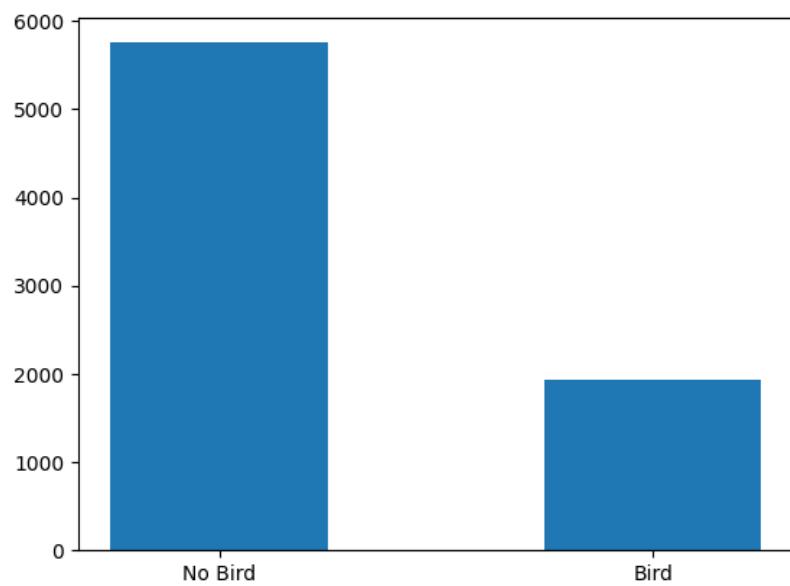


Figure 3.3: Dataset distribution for freefield1010

Another dataset used is the Crowdsourced dataset, UK ("warblrb10k") - 8,000 smartphone audio recordings from around the UK, crowdsourced by users of Warblr, the bird recognition app. The audio covers a wide distribution of UK locations and environments and includes weather noise, traffic noise, human speech, and even human bird imitations.

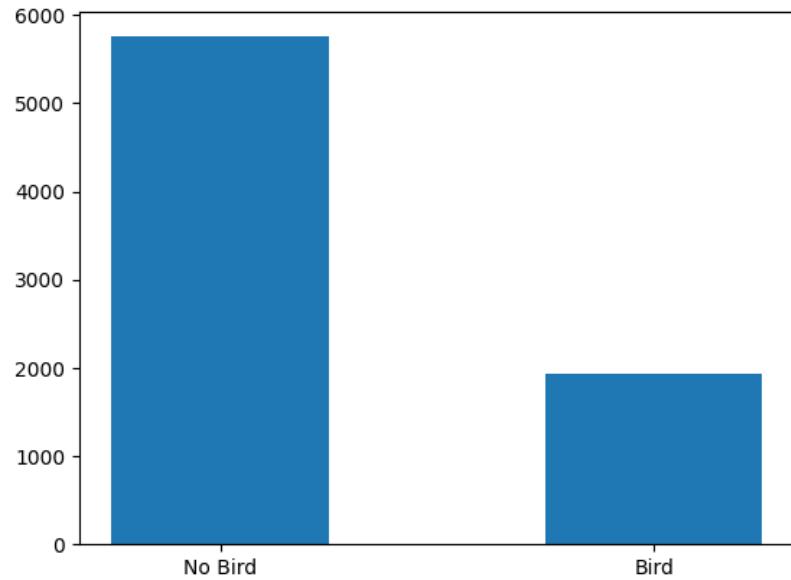


Figure 3.4: Dataset distribution for warblrb10k

Instead of using the datasets separately, we merged them to create a single dataset comprising 15,690 audio recordings—7,710 without bird sounds and 7,980 with bird sounds. The merged dataset was then partitioned into training, testing, and validation sets in an 80:10:10 ratio, ensuring a comprehensive and balanced dataset for model training and evaluation.

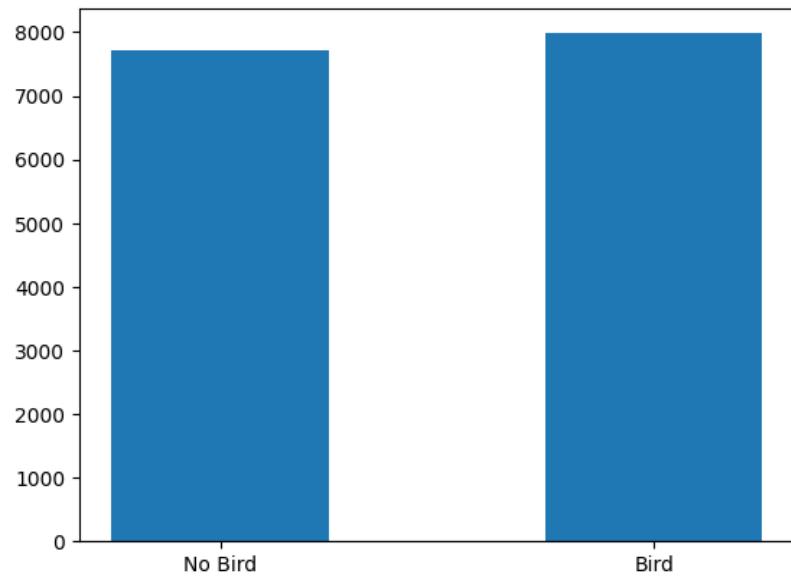


Figure 3.5: Distribution of the merged dataset

3.2.2 Data Transformation

To enhance the robustness and generalization of our bird sound detection model, we applied a series of transformations on both the raw audio signals and their corresponding Mel Spectrograms. These transformations include:

- **Add Gaussian Noise (Raw Audio):**

This transformation adds Gaussian noise to the raw audio signal. The noise is generated with a specified mean and standard deviation, and it is added to the audio signal with a certain probability. This helps in making the model robust to noisy environments.

$$\text{noise} = \mathcal{N}(\mu, \sigma^2) \quad (3.1)$$

- **Random Volume Scaling (Raw Audio):**

This transformation scales the volume of the raw audio signal by a random factor within a specified range. The scaling is applied with a certain probability, making the model resilient to variations in recording levels.

$$\text{audio} = \text{audio} \times \text{gain} \quad (3.2)$$

- **Time Stretch (Raw Audio):**

This transformation stretches or compresses the time axis of the raw audio signal by a random factor within a specified range. The transformed audio is then converted into a Mel Spectrogram, which helps the model to handle variations in speed.

$$\text{audio} = \text{TimeStretch}(\text{audio}, \text{rate}) \quad (3.3)$$

- **Frequency Masking (Mel Spectrogram):**

This transformation masks a portion of the frequency bins in the Mel Spectrogram. The maximum number of frequency bins to be masked is specified, and the masking is applied with a certain probability. It encourages the model to focus on different frequency components.

$$\text{mel_spec}[f : f + \text{max_mask}] = 0 \quad (3.4)$$

- **Time Masking (Mel Spectrogram):**

This transformation masks a portion of the time frames in the Mel Spectrogram. The maximum number of time frames to be masked is specified, and the masking is applied with a certain probability. This helps the model to focus on various temporal features.

$$\text{mel_spec}[:, t : t + \text{max_mask}] = 0 \quad (3.5)$$

- **Random Content Mixing (Mel Spectrogram):**

This transformation mixes the Mel Spectrogram with another randomly selected noise Mel Spectrogram. The mixing ratio is chosen randomly within a specified range, and the mixing is applied with a certain probability. This aids in making the model robust to background noises.

$$\text{mel_spec} = (1 - \text{mix_ratio}) \times \text{mel_spec} + \text{mix_ratio} \times \text{noise_spec} \quad (3.6)$$

- **Time Interval Dropout (Mel Spectrogram):**

This transformation randomly drops intervals of time frames in the Mel Spectrogram. The number of intervals and the width of each interval are chosen randomly within specified ranges, applied with a certain probability. This helps the model to cope with missing or corrupted audio segments.

$$\text{mel_spec}[:, t : t + \text{width}] = 0 \quad (3.7)$$

3.2.3 Feature Extraction

The Mel Spectrogram transforms raw audio signals into a time-frequency image that captures the essential spectral characteristics required for bird sound detection. As demonstrated by Lasseck et al.[11], this approach efficiently distinguishes recordings containing bird vocalizations from those without. For a comprehensive explanation of the computation, please refer to the Feature Extraction.

3.2.4 Model Architecture

The bird sound detection model utilizes a pre-trained ResNet-50 architecture as a feature extractor, which has been modified to process mel spectrogram representations of audio signals. This section provides a detailed explanation of the ResNet-50 architecture and the modifications applied to adapt it for bird sound detection.

3.2.4.1 ResNet-50: A Detailed Overview

ResNet-50, or Residual Network-50, is a deep convolutional neural network (CNN) designed to overcome the vanishing gradient problem that commonly arises in deep networks. It achieves this through the use of residual learning, where shortcut connections (skip connections) allow gradients to bypass certain layers, ensuring efficient training even in very deep architectures.

3.2.4.2 Architecture of ResNet-50

ResNet-50 consists of multiple convolutional layers arranged into residual blocks. These blocks introduce identity mappings that help the model learn transformations efficiently. The network follows a hierarchical structure where initial layers focus on low-level features such as edges and textures, while deeper layers extract high-level representations.

The architecture is composed of four main stages, each containing bottleneck residual blocks. Each block consists of three convolutional layers: a 1×1 convolution for dimensionality reduction, a 3×3 convolution for feature extraction, and another 1×1 convolution to restore dimensionality. By utilizing these bottleneck layers, ResNet-50 maintains computational efficiency while preserving representational power.

Another key component of ResNet-50 is batch normalization, which standardizes feature maps at each layer, leading to stable gradient flow and faster convergence. Additionally, ReLU activation functions introduce non-linearity, allowing the network to learn complex patterns in the data.

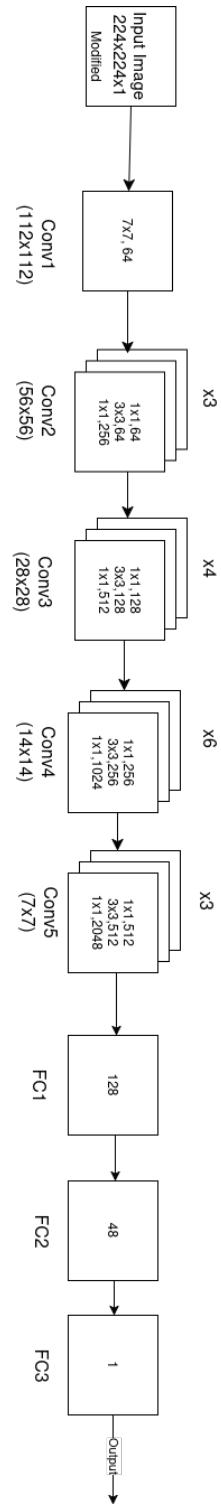


Figure 3.6: Modified ResNet-50 architecture for Bird Sound Detection.

3.2.4.3 Components of ResNet-50 Architecture

The ResNet-50 model consists of several essential components that contribute to its powerful feature extraction and classification capabilities:

- **Convolutional Layers:** These layers apply convolution operations to extract spatial features from input images. Early layers capture basic features (edges, textures), while deeper layers learn more abstract patterns.
- **Batch Normalization:** This technique normalizes activations at each layer to stabilize training, improve convergence, and prevent internal covariate shifts.
- **ReLU Activation:** A non-linear activation function that introduces sparsity and helps the network learn complex representations.
- **Residual Blocks:** The defining characteristic of ResNet, these blocks contain skip connections that allow gradients to bypass certain layers, mitigating the vanishing gradient problem.
- **Bottleneck Layers:** Each residual block consists of three convolutional layers: a 1×1 convolution to reduce dimensionality, a 3×3 convolution to extract features, and another 1×1 convolution to restore the original depth.
- **Global Average Pooling:** Reduces the dimensionality of feature maps before passing them to the final classification layer, making the network more efficient.
- **Fully Connected Layer:** The final layer of the network that performs classification based on extracted features.
- **Softmax/Sigmoid Activation:** Applies a probability distribution to the output for classification, with softmax used for multi-class classification and sigmoid used for binary classification.

3.2.4.4 Adaptation for Bird Sound Detection

While ResNet-50 is originally designed for image classification, it can be adapted for processing spectrograms, which are visual representations of sound frequencies over time. In this model, several modifications have been applied to tailor ResNet-50 for the bird sound detection task (see Figure 3.6):

- **Input Modification:** Since spectrograms are typically represented as grayscale images with a single channel, the original ResNet-50 architecture, which expects three-channel RGB images, has been adjusted to accept single-channel inputs. This allows the network to effectively process spectrogram data without requiring

redundant channels.

- **Feature Extraction:** The pre-trained ResNet-50 model, trained on large-scale image datasets, is used as a feature extractor. The earlier layers remain unchanged, as they capture fundamental patterns useful for various classification tasks. These convolutional layers are frozen, meaning their weights are not updated during training. This strategy ensures that the model benefits from pre-learned visual features while focusing its learning on the final classification task.
- **Modified Classification Layer:** The original fully connected classification layer of ResNet-50 has been replaced with a custom classification head optimized for binary classification (presence or absence of bird sound). This new layer stack introduces additional non-linearity and dropout mechanisms to prevent overfitting, ultimately improving the model’s ability to generalize to new data.

3.2.4.5 Training and Generalization

By freezing the early convolutional layers and training only the newly added classification layers, the model efficiently learns to distinguish bird sounds from other environmental noises. This transfer learning approach significantly reduces the training time and the amount of labeled data required, making the model robust in real-world scenarios.

The final model is designed to process mel spectrogram inputs and produce a probability score indicating the presence of bird sounds. The use of residual learning allows for deep feature extraction, making the model highly effective in distinguishing between subtle variations in audio patterns.

3.2.5 Training Process

The training process for the bird sound detection model was carefully designed to ensure optimal learning and generalization. The process incorporated several key components and strategies:

- **Training Parameters:**

- Number of epochs: 20
- Batch size: 32
- Initial learning rate: 0.001
- Weight decay: 0.0001
- **Loss Function:** Binary Cross-Entropy Loss (BCE) was chosen as the loss function due to its effectiveness in binary classification tasks. For input predictions \hat{y} and true labels y , the BCE loss is calculated as:

$$L_{BCE} = -\frac{1}{N} \sum_{i=1}^N [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)] \quad (3.8)$$

where N is the number of data samples.

- **Optimization:**
 - The Adam optimizer was employed with the following parameters:
 - * Learning rate: 0.001
 - * Weight decay: 0.0001
 - Learning rate scheduling was implemented using ReduceLROnPlateau with:
 - * Reduction factor: 0.1
 - * Patience: 3 epochs
- **Monitoring and Validation:**
 - Training and validation losses were monitored each epoch
 - Learning rate adjustments were made based on validation loss plateaus
 - Early stopping was implemented to prevent overfitting
 - Model performance was evaluated on the test set after training completion

The key hyperparameters used for training are summarized in Table 3.2.

3.2.6 Evaluation

To assess the effectiveness of our model in distinguishing between the target classes, we analyzed its performance using the Receiver Operating Characteristic (ROC) curve. Following the benchmark established by Lasseck et al.[11] for the freefield1010 and warblrb10k datasets, the Area Under the Curve (AUC) metric was chosen as the primary

Table 3.1: Training Hyperparameters

Hyperparameter	Value
Number of Epochs	20
Batch Size	32
Initial Learning Rate	0.001
Weight Decay	0.0001
Loss Function	Binary Cross-Entropy (BCE)
Optimizer	Adam
Learning Rate Scheduling	ReduceLROnPlateau
Reduction Factor	0.1
Patience	3 epochs
Early Stopping	Patience = 3 epochs

evaluation criterion.

The ROC curve provides a comprehensive visualization of the model's discriminative ability by plotting the True Positive Rate (TPR) against the False Positive Rate (FPR) at various classification thresholds. For our binary classification task:

- **True Positive Rate (Sensitivity):**

$$TPR = \frac{TP}{TP + FN} \quad (3.9)$$

Measures the proportion of actual bird sounds correctly identified.

- **False Positive Rate (1 - Specificity):**

$$FPR = \frac{FP}{FP + TN} \quad (3.10)$$

Represents the proportion of non-bird sounds incorrectly classified as bird sounds.

‘’ The Area Under the ROC Curve (AUC) provides a single scalar value representing the model's overall classification performance. An AUC of 1.0 represents perfect classification, while 0.5 indicates random chance performance.

The ROC curve analysis reveals several important aspects of our model's performance:

- **Threshold Independence:** The ROC curve demonstrates the model's perfor-

mance across all possible classification thresholds, providing a more robust evaluation than single-threshold metrics.

- **Class Imbalance Handling:** Unlike accuracy, the ROC curve remains effective even when dealing with imbalanced datasets, making it particularly suitable for real-world bird sound detection scenarios.
- **Operating Point Selection:** The curve enables the selection of optimal decision thresholds based on specific requirements for sensitivity versus specificity in deployment scenarios.

The high AUC score indicates that our model effectively separates bird sounds from background noise and other environmental sounds, making it reliable for real-world applications in bird sound detection.

3.3 Working Mechanism for Bird Species Classification

The working mechanism for bird identification from audio utilizes the dataset compiled from Xeno Canto, housing a diverse collection of avian vocalizations. Our methodology revolves around the utilization of EfficientNet, a state-of-the-art convolutional neural network architecture known for its balance between accuracy and efficiency.

The primary objective of this study is to achieve high levels of accuracy in identifying a broad spectrum of 41 distinct bird species. Here lies the detailed explanation of the methodology, from the working mechanism to model training.

3.3.1 Data Collection

For the bird species classification model, we scraped audio data from Xeno-canto (*xeno-canto.org*), a global repository of bird vocalizations contributed by ornithologists and birding enthusiasts. Since the dataset was highly imbalanced, we applied a series of pre-processing and augmentation techniques to ensure a more uniform distribution across bird species.




Figure 3.7: Block diagram for the working mechanism for Bird Species Identification

3.3.2 Data Augmentation

Initial analysis of the collected data revealed significant class imbalance, with some species having over 400 recordings while others had fewer than 50. Additionally, the audio recordings varied in length. To address these issues:




Figure 3.8: Dataset before performing augmentation

- **Audio Clipping:**

- Recordings were clipped into segments of 10 seconds each.
- Clips shorter than 5 seconds were discarded to avoid blank audio segments and insufficient data.
- Clips between 5 and 10 seconds were padded with silence at the end to standardize their length to 10 seconds.

- **Augmentation Techniques:** To balance the dataset, augmentation techniques such as time stretching, phase shifting, and noise addition were applied. The parameters were varied to ensure diversity in the augmented data.

- Each bird class was augmented to contain exactly 500 audio clips.

3.3.3 Data Preprocessing

The audio recordings were transformed into Mel Spectrograms for further analysis.

- **Conversion Details:**

- Audio files were converted using a sample rate of 32,000 Hz.
- The spectrograms were generated with a Hanning window and 48 Mel bands.

- **Dataset Expansion:**

- Each audio file was converted into corresponding Mel Spectrogram images




Figure 3.9: Dataset after performing augmentation

for CNN based models and MFCC vectors for LSTM based model.

- The resulting dataset was divided into training, validation, and testing sets for model development.

3.3.4 Feature Extraction

The Mel Spectrogram was employed to transform raw audio recordings into a two-dimensional time-frequency representation, effectively converting sound into an image-like format for processing by convolutional neural networks. This method captures both the temporal and spectral nuances of bird calls, allowing the classifier to discern distinguishing features among various species. As demonstrated in the work of Sevilla et al.[5], such time-frequency representations enhance classification accuracy by emphasizing key frequency components inherent in bird vocalizations.

For a more detailed discussion on the Mel Spectrograms, please refer to the 3.4.

3.3.5 Model Architecture

The bird call classification model is based on EfficientNetB3, a deep convolutional neural network (CNN) designed for optimal efficiency and accuracy. EfficientNetB3

achieves this through **compound scaling**, which uniformly scales network depth, width, and resolution to maximize performance while minimizing computational cost.

3.3.5.1 EfficientNetB3

EfficientNetB3 is part of the EfficientNet family, which employs a balance between deeper layers, wider networks, and higher input resolutions. This architecture integrates several enhancements, including **Mobile Inverted Bottleneck Convolution (MBConv) blocks**, **squeeze-and-excitation (SE) modules**, and the **Swish activation function**. These components contribute to better feature extraction and gradient flow, ensuring robust classification performance.

i Architecture of EfficientNetB3

EfficientNetB3 follows a hierarchical feature extraction process, where early layers detect fundamental patterns (edges, textures), while deeper layers extract more complex features. The architecture consists of multiple MBConv blocks that improve computational efficiency by using **depthwise separable convolutions** and **linear bottlenecks**. Additionally, SE modules dynamically recalibrate channel-wise feature responses, enhancing the model's ability to focus on important frequency patterns in the mel spectrogram.

The key architectural components of EfficientNetB3 are summarized below:

- **Convolutional Layers:** Extract spatial features from input spectrograms.
Early layers capture basic patterns (edges, textures), while deeper layers learn abstract audio representations.
- **MBConv Blocks:** Mobile Inverted Bottleneck Convolution blocks enhance efficiency by using depthwise separable convolutions and linear bottlenecks.
- **Squeeze-and-Excitation (SE) Modules:** Improve feature recalibration by dynamically adjusting channel-wise feature responses, allowing the network to focus on relevant frequency patterns.
- **Swish Activation:** A smooth, non-monotonic activation function that enhances gradient flow and optimization stability.
- **Batch Normalization:** Standardizes activations at each layer to stabilize




Figure 3.10: Modified EfficientNetB3 architecture for Bird Species Classification.

training and improve convergence.

- **Global Average Pooling:** Reduces dimensionality of feature maps before

passing them to the classification layer.

- **Dropout Layer:** A dropout rate of 0.3 is applied to prevent overfitting by randomly deactivating neurons.
- **Fully Connected Layer:** Converts extracted features into classification outputs.
- **Softmax Activation:** Generates probability distributions for multi-class classification.

ii) **Adaptation for Bird Call Classification**

While EfficientNetB3 is originally designed for image classification, it has been adapted for bird call classification by processing mel spectrograms as input. The key modifications include:

- **Input Modification:** The model accepts spectrogram inputs with dimensions $128 \times 128 \times 3$, where three channels correspond to different spectrogram augmentations or frequency bands.
- **Feature Extraction:** The EfficientNetB3 layers pre-trained on ImageNet are used as a feature extractor. The convolutional layers are **frozen** to retain general visual features while training the classification head on bird call data.
- **Custom Classification Head:** The original fully connected layers are replaced with a new classification head consisting of a **Flatten** layer, **Dropout** layer (0.3), and a **Dense** layer. A softmax activation function is applied to generate class probabilities.

iii) **Training and Generalization**

The model is trained using a transfer learning approach, where only the classification layers are updated while the pre-trained EfficientNetB3 layers remain frozen. This reduces training time and improves generalization.

iv) **Training Process**

– **Initial Training Parameters:**

- * Number of epochs: 30
- * Batch size: 4
- * Image size: $128 \times 128 \times 3$
- * Initial learning rate: 0.0001

- **Loss Function:** Categorical Cross-Entropy(CCE) was chosen as the loss function since it is well-suited for multi-class classification tasks. The Categorical Cross-Entropy is calculated as:

$$L_{CCE} = -\frac{1}{N} \sum_{i=1}^N \sum_{j=1}^C y_{i,j} \log(\hat{y}_{i,j}) \quad (3.11)$$

where N is Number of data samples and C is number of Classes.

- **Optimization:**

- * The Adam optimizer was employed with the following parameters:
 - Learning rate: 0.0001
- * Learning rate scheduling was implemented using an exponential decay function after the first five epochs.

- **Regularization:**

- * Dropout (0.3) was applied to prevent overfitting.
- * L2 regularization ($\lambda = 1e^{-4}$) was used for additional weight regularization.

- **Monitoring and Validation:**

- * Training and validation losses were monitored each epoch.
- * Learning rate adjustments were made using an exponential decay function.
- * Early stopping was implemented with a patience of three epochs to prevent overfitting.
- * Model checkpointing was applied to save the best-performing model based on validation loss.

The key hyperparameters used for training are summarized in Table 3.2.

The final model efficiently classifies bird calls by leveraging deep feature extraction from EfficientNetB3 while incorporating a custom classification head optimized for spectrogram data. Through transfer learning and regularization techniques, the model generalizes well to new audio samples, achieving high classification accuracy.

Table 3.2: Initial Training Hyperparameters

Hyperparameter	Value
Number of Epochs	30
Batch Size	4
Input Image Size	$128 \times 128 \times 3$
Initial Learning Rate	0.0001
Loss Function	Categorical Cross-Entropy
Optimizer	Adam
Learning Rate Scheduling	Exponential Decay (after 5 epochs)
Dropout Rate	0.3
L2 Regularization	$1e^{-4}$
Early Stopping	Patience = 3 epochs

3.3.5.2 CNN+LSTM Model

The CNN+LSTM architecture is designed for bird call classification using MFCC (Mel-Frequency Cepstral Coefficients) matrices instead of mel spectrogram images. Unlike EfficientNetB3, which processes image representations of audio, this model extracts both spatial and temporal features by combining convolutional layers with recurrent layers.

3.3.5.3 Architecture of CNN+LSTM

This model follows a two-stage feature extraction process:

- **Convolutional Feature Extraction:** Convolutional layers extract local frequency features from MFCC matrices, capturing patterns in short time frames.
- **Temporal Modeling with LSTM:** A Long Short-Term Memory (LSTM) layer processes sequential dependencies in the extracted features, enhancing temporal understanding.

The key architectural components of the CNN+LSTM model are as follows:

- **Convolutional Layers:** Three convolutional layers extract spatial features from MFCC input. Each convolutional block includes:
 - **Conv2D:** Filters of sizes (3×3) with ReLU activation.

- **Batch Normalization:** Normalizes activations for stable training.
- **MaxPooling2D:** Reduces feature map dimensions by pooling along the time axis.
- **Time-Distributed Flattening:** Flattens convolutional feature maps while preserving temporal dimensions for LSTM processing.
- **LSTM Layer:** A 256-unit LSTM layer captures long-range temporal dependencies in bird call sequences.
- **Fully Connected Layers:** Two dense layers process the LSTM output:
 - A 64-unit dense layer with ReLU activation and dropout (0.3) for regularization.
 - An output dense layer with softmax activation for classification.

3.3.5.4 Adaptation for Bird Call Classification

The CNN+LSTM model is tailored to classify bird calls based on MFCC matrices instead of image-based spectrograms. The key modifications include:

- **Input Modification:** The model accepts MFCC matrices with dimensions $(13 \times 626 \times 1)$, where 13 represents MFCC coefficients, and 626 represents time frames.
- **Sequential Feature Extraction:** CNN layers extract spatial frequency information, while the LSTM layer captures sequential dependencies in bird calls.
- **Custom Classification Head:** Includes fully connected layers with dropout to reduce overfitting.

3.3.5.5 Training and Generalization

The model is trained using a combination of early stopping, learning rate scheduling, and model checkpointing to ensure optimal performance.

3.3.5.6 Training Process

- **Initial Training Parameters:**
 - Number of epochs: 30

- Batch size: 16
- Learning rate: 0.0001
- Loss function: Categorical Cross-Entropy
- **Optimization:**
 - Adam optimizer with an initial learning rate of 0.0001.
 - Learning rate scheduling using an exponential decay function after the first five epochs.
- **Regularization:**
 - Dropout (0.3) applied to prevent overfitting.
- **Monitoring and Validation:**
 - Training and validation losses monitored per epoch.
 - Early stopping with a patience of 3 epochs.
 - Model checkpointing to save the best-performing model based on validation loss.

The key hyperparameters used for training are summarized in Table 3.3.

Table 3.3: Initial Training Hyperparameters for CNN+LSTM

Hyperparameter	Value
Number of Epochs	30
Batch Size	16
Input Shape	(13, 626, 1)
Learning Rate	0.0001
Loss Function	Categorical Cross-Entropy
Optimizer	Adam
Learning Rate Scheduling	Exponential Decay (after 5 epochs)
Dropout Rate	0.3
Early Stopping	Patience = 3 epochs

The final CNN+LSTM model effectively classifies bird calls by integrating deep convolutional feature extraction with temporal modeling, ensuring high accuracy while leveraging MFCC representations.

3.3.6 Evaluation Metrics

To assess the performance of the proposed models, the following evaluation metrics were used. These metrics provide a comprehensive understanding of the models classi-

fication ability by analyzing both correct predictions and class-wise performance.

- **Accuracy:** Accuracy measures the overall correctness of the model by calculating the ratio of correctly classified samples to the total number of samples.

$$Accuracy = \frac{TP + TN}{TP + TN + FP + FN} \quad (3.12)$$

- **Recall:** Recall (also known as sensitivity or true positive rate) measures the model's ability to correctly identify positive instances.

$$Recall = \frac{TP}{TP + FN} \quad (3.13)$$

- **Precision:** Precision quantifies the proportion of correctly predicted positive instances out of all predicted positive instances.

$$Precision = \frac{TP}{TP + FP} \quad (3.14)$$

- **F1-score:** The F1-score is the harmonic mean of precision and recall, providing a balanced measure when there is an uneven class distribution.

$$F1\text{-score} = 2 \times \frac{Precision \times Recall}{Precision + Recall} \quad (3.15)$$

3.4 Feature Extraction

Feature extraction was a critical step where the audio clips were transformed into a format that were fed into machine learning models. We used two primary techniques for feature extraction: Mel Spectrograms and Mel-Frequency Cepstral Coefficients (MFCC).

3.4.1 Mel Spectrogram

A spectrogram is a visual representation of the spectrum of frequencies in a sound signal as they vary with time. It is generated by applying the Short-Time Fourier Transform (STFT) to the audio signal, followed by mapping the frequencies to the Mel scale, which

better represents human auditory perception. This transformation provides insight into how the frequency content of the signal changes over time.

- 1. Short-Time Fourier Transform (STFT):** The STFT is computed by dividing the signal into overlapping frames, applying a window function, and then performing the Fourier Transform on each frame.

$$X(n) = \sum_{m=0}^{N-1} x(m) \cdot w(n-m) \quad (3.16)$$

where $x(m)$ is the audio signal and $w(n)$ is the window function.

- 2. Mel Scale Transformation:** The frequency axis of the spectrogram is then mapped to the Mel scale using triangular filter banks.

$$M(f) = 2595 \log_{10} \left(1 + \frac{f}{700} \right) \quad (3.17)$$

where $M(f)$ is the Mel frequency, and f is the frequency in Hz.

- 3. Log Scaling and Normalization:** Convert Mel spectrogram to a logarithmic scale:

$$S_{\log}(m, n) = \log(S_{\text{mel}}(m, n)) \quad (3.18)$$

where $S_{\log}(m, n)$ is the log-scaled Mel spectrogram, and $S_{\text{mel}}(m, n)$ is the Mel spectrogram before log transformation.




Figure 3.11: Mel Spectrogram image for a sample audio

3.4.2 Mel-Frequency Cepstral Coefficients (MFCC)

MFCCs are coefficients that collectively describe the short-term power spectrum of a sound signal. The process of obtaining MFCCs involves several steps:




Figure 3.12: Feature Extraction Using Spectrogram and MFCC

- Framing:** Divide the audio signal into short frames.
- Discrete Fourier Transform (DFT):** Convert each frame to the frequency domain.

$$X(k) = \sum_{n=0}^{N-1} x(n) \cdot e^{-j \frac{2\pi}{N} kn} \quad (3.19)$$

- Log Function:** Apply a logarithm to the amplitude spectrum.

$$S_{\log}(k) = \log(|X(k)|) \quad (3.20)$$

- Mel-Scaling:** Map the frequencies to the Mel scale, which better represents how humans perceive sound.

$$f_{\text{mel}} = 2595 \cdot \log_{10}\left(1 + \frac{f}{700}\right) \quad (3.21)$$

- Discrete Cosine Transform (DCT):** Convert the Mel spectrum to the cepstral domain, yielding the MFCC features.

$$C(n) = \sum_{k=0}^{K-1} S_{\text{mel}}(k) \cdot \cos\left(\frac{\pi n(k+0.5)}{K}\right) \quad (3.22)$$




Figure 3.13: MFCC image for a sample audio

3.5 Hyperparameter Optimization using Genetic Algorithm

Genetic algorithms are a powerful method for optimizing hyperparameters in machine learning models. Genetic Algorithm have proven to significantly improve the performance metrics of the CNN model instead of using hand tuned approach for hyperparameters. This section outlines the steps involved in using GAs for hyperparameter optimization[12].

- **Encoding the Hyperparameters**

- Hyperparameters are represented as a chromosome, where each hyperparameter is a gene in the chromosome.
- For example, in a neural network, a chromosome might include genes for the learning rate, number of layers, number of neurons per layer, and activation functions.

- **Initial Population**

- An initial population of chromosomes is generated randomly, with each chromosome representing a different set of hyperparameters.

- **Fitness Function**

- A fitness function is defined to evaluate the performance of each set of hyperparameters.
- This typically involves training the model with the given hyperparameters and measuring its performance on a validation set.

- **Selection**

- Selection involves choosing the best-performing chromosomes to serve as parents for the next generation.
- Various selection methods can be employed, such as tournament selection, roulette wheel selection, or rank-based selection.

- **Crossover (Recombination)**

- Crossover combines pairs of parent chromosomes to produce offspring for the next generation.
- This is done by swapping segments of parent chromosomes to create new chromosomes, thereby combining features of both parents.

- **Mutation**

- Mutation introduces random changes to some of the genes in the offspring chromosomes.
- This helps maintain genetic diversity in the population and allows the algorithm to explore a broader search space.

- **Replacement**

- The current population is partially or entirely replaced with the new generation of chromosomes, ensuring that better solutions are carried forward while allowing for exploration of new possibilities.

- **Termination**

- The process of selection, crossover, mutation, and replacement is repeated until a termination criterion is met.
- This could be a set number of generations, convergence of fitness scores, or achieving a satisfactory performance level.

- **Best Solution**

- The best chromosome at the end of the process represents the optimal or near-optimal set of hyperparameters for the model.

CHAPTER 4

RESULTS AND ANALYSIS

In the course of our project on Bird Species Classification from Audio, we have accomplished several key tasks to progress our analysis and model development:

4.1 Bird Sound Detection

Figure 4.3 shows the training progress over 15 epochs. The model shows consistent improvement in both accuracy and loss metrics during training and validation phases.

The training process exhibited steady improvement, with the model achieving final metrics of:

- Training accuracy: **84.20%** and validation accuracy: **84.10%**
- Training loss: **0.394** and validation loss: **0.412**
- Training F1-score: **0.585** and validation F1-score: **0.615**

The close alignment between training and validation metrics indicates good generalization without overfitting. The loss curves show consistent convergence, while accuracy and F1-scores demonstrate steady improvement throughout the training process.

The confusion matrix in Figure 4.4 shows the classification performance on the test dataset. From the total of **1621** samples, the model correctly classified **621** samples as non-bird sounds (True Negatives) and 680 samples as bird sounds (True Positives). There were 107 false positives and **160** false negatives. The model performance metrics are summarized in Table 4.1.

Table 4.1: Performance Metrics for Bird Sound Detection Model

Metric	Value (%)
Accuracy	82.97
F1-score	83.59
Precision	80.95
Recall	86.40

Figure 4.5 shows the Receiver Operating Characteristic (ROC) curve for the bird sound



Figure 4.1: Accuracy curves




Figure 4.2: Loss curves

Figure 4.3: Training and validation curves for Bird Sound Detection Model

detection model. The curve illustrates the model's ability to discriminate between bird and non-bird sounds across different classification thresholds.




Figure 4.4: Confusion Matrix for Bird Sound Detection Model




Figure 4.5: ROC Curve of the detection model

The Area Under the Curve (AUC) score of **0.911** indicates excellent discriminative ability, as it is significantly higher than the random chance baseline of **0.5**.

4.2 Bird Species Classification

The below figures show the outcomes of Bird Species Classification based on two different architectures i.e CNN-LSTM and EfficientNetB3, evaluated on our dataset of 41 bird species.

4.2.1 Classification using CNN-LSTM

The CNN-LSTM model achieved an accuracy of **74.73%** with a cross-entropy loss of **1.0171** across all 41 bird species. The confusion matrix shows stronger diagonal elements indicating good classification ability for most species, though there is some confusion between similar-sounding birds. The training and validation curves show the




Figure 4.6: CNN-LSTM Model Evaluation

model achieved a training accuracy of **80.76%** and validation accuracy of **74.73%**, with corresponding losses of **0.8025** and **1.0171** respectively.




Figure 4.7: CNN-LSTM Confusion Matrix

The confusion matrix visualizes the model's predictions across all 41 classes, with darker diagonal elements indicating correct classifications and lighter off-diagonal elements showing misclassifications between species.




Figure 4.8: CNN-LSTM Classification Report

The classification report shows detailed metrics for each species, with an average precision of **73%**, recall of **75%**, and F1-score of **73%** across all classes.

4.2.2 Classification using EfficientNetB3

The EfficientNetB3 model demonstrated superior performance with an accuracy of **90.73%** with a cross-entropy loss of **0.4154**, significantly outperforming the CNN-LSTM architecture. The confusion matrix reveals excellent discrimination ability across most species, with minimal misclassifications between different bird calls.




Figure 4.9: EfficientNetB3 Model Evaluation

The training and validation curves demonstrate strong performance with a training accuracy of **93.32%** and validation accuracy of **91.61%**, while maintaining low loss values of **0.2336** and **0.3706** respectively.




Figure 4.10: EfficientNetB3 Confusion Matrix

The confusion matrix shows strong diagonal elements indicating excellent classification accuracy across most species, with minimal confusion between different bird calls.




Figure 4.11: EfficientNetB3 Classification Report

The classification report shows detailed metrics for each species, with an average precision, recall and f1-score of **91%** across all classes.

4.3 Hyperparameter Optimization Using Genetic Algorithm

To optimize the performance of our model, we used genetic algorithm for hyperparameter tuning. This approach helped us systematically explore and identify the most effective hyperparameter settings for our model, including network architecture parameters like learning rates, batch sizes, images sizes, dropout and others. The genetic algorithm enabled a more efficient and comprehensive search for optimal configurations.

4.3.1 Hyperparameters for CNN-LSTM

For the CNN-LSTM model, GA was applied to tune hyperparameters that impact both convolutional feature extraction and model efficiency. The key hyperparameters optimized includes Batch Size, Learning Rate, Number of Epochs, Dropout Rate, LSTM Units and Dense Units. The genetic algorithm was executed over multiple generations, evolving hyperparameter configurations towards optimal performance.

4.3.1.1 Hyperparameter Search Space

The genetic algorithm optimized the following hyperparameters:

Hyperparameter	Search Values
Batch Size	{16, 32, 64}
Learning Rate	{1e-4, 1e-4, 1e-3}
Epochs	{20, 30, 40}
Dropout Rate	{0.3, 0.4, 0.5}
LSTM Units	{64, 128, 256}
Dense Units	{64, 128, 256}

Table 4.2: Hyperparameter search space for CNN-LSTM

4.3.1.2 Genetic Algorithm Process

The GA followed an evolutionary approach, iterating over multiple generations to refine hyperparameters. The process consisted of:

- **Population Size:** 3 (each generation consists of 3 hyperparameter sets)
- **Generations:** 2 (iterative evolution of hyperparameters)

- **Mutation Rate:** 10% (randomly modifying one parameter)
- **Crossover:** Parents were selected based on validation accuracy, and offspring were created by combining their hyperparameters.

4.3.1.3 Results

Over two generations, the algorithm refined hyperparameters, improving model accuracy and minimizing overfitting. The following images illustrate training and validation performance trends for different hyperparameter configurations across generations.

The Following Graphs shows the training and validation accuracies and loss for the respective hyperparameters used for CNN-LSTM.




Figure 4.12: Accuracy and Loss for CNN-LSTM in generation 1




Figure 4.13: Accuracy and Loss for CNN-LSTM in generation 2




Figure 4.14: Accuracy and Loss for CNN-LSTM in generation 3




Figure 4.15: Accuracy and Loss for CNN-LSTM in generation 4




Figure 4.16: Accuracy and Loss for CNN-LSTM in generation 5

The following table demonstrates different hyperparameter configurations and their corresponding test accuracies.

SN	Batch	LR	Epochs	Dropout	LSTM Units	Dense Units	Test Accuracy (%)
1	64	0.001	40	0.5	128	64	7.61
2	16	0.001	30	0.3	256	64	25.76
3	16	0.0001	20	0.3	256	64	75.17
4	16	0.0001	20	0.4	256	64	65.80
5	16	0.001	30	0.3	65	64	11.61

Table 4.3: Hyperparameter configurations and test accuracies for CNN-LSTM

4.3.1.4 Conclusion

The best model achieved a Test Accuracy of **75.17%** with a **Batch Size of 16, Learning Rate of 0.0001, Epochs = 20, Dropout Rate of 0.3, LSTM Units = 256, and Dense Units = 64**. These results highlight the effectiveness of the Genetic Algorithm in systematically optimizing hyperparameters, leading to improved learning efficiency and minimized overfitting.

4.3.2 Hyperparameters for Efficient-Net

For EfficientNet, GA was used to fine-tune hyperparameters that directly affect feature extraction and classification efficiency. The optimized parameters includes Batch Size, Learning Rate, Number of Epochs, Dropout Rate, Regularization and Dense Units.

4.3.2.1 Hyperparameter Search Space

The genetic algorithm optimized the following hyperparameters:

Hyperparameter	Search Values
Batch Size	{4, 8, 16}
Image Size	{(128, 128), (128, 128), (224, 224)}
Learning Rate	{1e-4, 1e-4, 1e-3}
Epochs	{20, 20, 20}
Dropout Rate	{0.3, 0.2, 0.4}
L2 Regularization	{1e-4, 1e-4, 1e-3}

Table 4.4: Hyperparameter search space for the EfficientNetB3

4.3.2.2 Genetic Algorithm Process

The GA followed an evolutionary approach, iterating over multiple generations to refine hyperparameters. The process consisted of:

- **Population Size:** 3 (each generation consists of 3 hyperparameter sets)
- **Generations:** 2 (iterative evolution of hyperparameters)
- **Mutation Rate:** 10% (randomly modifying one parameter)

- **Crossover:** Parents were selected based on validation accuracy, and offspring were created by combining their hyperparameters.

4.3.2.3 Results

Over three generations, the algorithm refined hyperparameters, leading to improved validation accuracy. The following figures illustrate training and validation performance trends for different hyperparameter configurations across generations.

The Following Graphs shows the training and validation accuracies and loss for the respective hyperparameters used for EfficientNetB3.




Figure 4.17: Accuracy and Loss for EfficientNetB3 in generation 1




Figure 4.18: Accuracy and Loss for EfficientNetB3 in generation 2




Figure 4.19: Accuracy and Loss for EfficientNetB3 in generation 3




Figure 4.20: Accuracy and Loss for EfficientNetB3 in generation 4




Figure 4.21: Accuracy and Loss for EfficientNetB3 in generation 5

The following table demonstrates different hyperparameter configurations and their corresponding test accuracies.

SN	Batch	Img Size	LR	Epochs	Dropout	L2	Test Accuracy (%)
1	8	224x224	0.0001	20	0.2	1e-4	93.37
2	8	128x128	0.001	20	0.4	1e-4	92.54
3	16	128x128	0.0001	20	0.3	1e-4	90.68
4	8	128x128	0.0001	20	0.2	1e-4	91.02
5	8	128x128	0.001	20	0.2	1e-4	92.54

Table 4.5: Hyperparameter configurations for Efficient-NetB3

4.3.2.4 Conclusion

The best model achieved a Test Accuracy of **93.37%** with a **Batch Size of 8, Image Size of 224x224, Learning Rate of 0.0001, Epochs = 20, Dropout Rate of 0.2, and L2 Regularization = 1e-4**. These results highlight the effectiveness of the Genetic Algorithm in systematically optimizing hyperparameters, leading to improved model performance and reduced overfitting.

4.4 Mobile Application

The mobile application allows the user to record the audio, visualize it with waveforms during recording, before sending the actual audio data to the server for processing. The user can also use the application to add their current location to the map. The homepage of the application allows the user to view the birds along with the audio that are available in FeatherFind's database. The map page in FeatherFind allows the user to view all the birds that have been located using the application and view their location in the maps embedded in the application alone with the detailed description of the birds.

For the audio recognition the user can record the audio in the application and as the audio is being recorded the app visualizes the audio in appropriate waveform. When the user stops recording the audio, the user is prompted with a confirmation dialog, which can be used to process the audio and recognize it.




Figure 4.22: Home Page of FeatherFind.




Figure 4.23: Mapped Birds in FeatherFind.




Figure 4.24: Recording using FeatherFind.




Figure 4.25: Confirmation for Audio Recognition.




Figure 4.26: Bird Species Identification using FeatherFind.

CHAPTER 5

DISCUSSION

The FeatherFind project has made significant progress in developing an advanced species identification using audio recordings. Some the key achievements of the project include:

5.1 Achievements

- Comprehensive Data Collection: The project successfully comiled an extensive dataset of d sounds by scraping Xeno-Canto, ensuring diversity in species representaion.
- Advanced Data Processing Techniques: The project utilized advancd audio processing techniques along with data augmentaion(time stretching,phase shifting, gradient noiise addition), were implemented to balance dataset and enhance model robustness.
- Feature Extraction and Model Training: Spectrograms and Mel-Frequency Cepstral Coefficients (MFCCs) were employed for feature extraction. EfficientNetB3 architecture was utilized for bird species classification, achieving an accuracy of 89.36% on the test dataset. Avolutional Neural Network (CNN) combined with Long ShoTerm Memory (LSTM) networks was used to capture both spatial ad temporal features.
- Genetic Algorithm for Hyperparameter Optimization: The model's performance was fine-tuned using genetic algorithms, ensuring optimal parameter selection for improved accuracy.
- Bird Sound Detection Model: A preliminary model was trained to verify the presence of bird sounds before classification, using the InceptionV3 network.AUC score of 83% and an accuracy of 87.28% was achieved.
- Deployment and Mobile App Development: The classification model was deployed on Huggingface Spaces for seamless API integration.A mobile application was developed to allow users to record and analyze bird sounds in real-time while tagging locations using GPS integration.

5.2 Limitations

- Dataset limitations : Since the dataset lacks sufficient variation in bird species, environmental noise, and recording conditions, FeatherFind struggles with real-world scenarios. Additionally, because certain bird species are overrepresented, the model produces biased predictions, making it less reliable for species with fewer training samples.
- Multiple bird sound : Since the system is optimized to recognize only a single dominant bird call in an audio recording, it fails to accurately identify multiple birds vocalizing simultaneously. When overlapping sounds occur, the model either misclassifies them or ignores background calls, limiting its ability to handle complex audio environments.
-

5.3 Future Improvements

REFERENCES

- [1] I. O. Union, “Ioc world bird list,” <https://www.worldbirdnames.org/new/updates/>, 2024, accessed: 2024-06-04.
- [2] H. Nature, “National red list of nepal’s birds,” <https://www.himalayannature.org/works/projects/national-red-list-of-nepals-birds/>, 2024, accessed: 2024-06-03.
- [3] C. Inskip, H. S. Baral, T. Inskip, A. P. Khatiwada, M. P. Khatiwada, L. P. Poudyal, and R. Amin, “Nepal’s national red list of birds,” *Journal of Threatened Taxa*, vol. 9, no. 1, pp. 9700–9722, 2017.
- [4] R. Gautam, B. Khatiwada, B. P. Subedi, N. Duwal, and K. C. Dahal, “Audio classifier for automatic identification of endangered bird species of nepal,” 2023.
- [5] A. Sevilla and H. Glotin, “Audio bird classification with inception-v4 extended with time and time-frequency attention mechanisms.” *CLEF (Working Notes)*, vol. 1866, pp. 1–8, 2017.
- [6] B. Chandu, A. Munikoti, K. S. Murthy, G. Murthy, and C. Nagaraj, “Automated bird species identification using audio signal processing and neural networks,” in *2020 International Conference on Artificial Intelligence and Signal Processing (AISP)*. IEEE, 2020, pp. 1–5.
- [7] L. Nanni, G. Maguolo, S. Brahnam, and M. Paci, “An ensemble of convolutional neural networks for audio classification,” *Applied Sciences*, vol. 11, no. 13, p. 5796, 2021.
- [8] H. Wang, Y. Xu, Y. Yu, Y. Lin, and J. Ran, “An efficient model for a vast number of bird species identification based on acoustic features,” *Animals*, vol. 12, no. 18, p. 2434, 2022.
- [9] R. Pahuja and A. Kumar, “Sound-spectrogram based automatic bird species recognition using mlp classifier,” *Applied Acoustics*, vol. 180, p. 108077, 2021.
- [10] S. Carvalho and E. F. Gomes, “Automatic classification of bird sounds: using mfcc and mel spectrogram features with deep learning,” *Vietnam Journal of Computer Science*, vol. 10, no. 01, pp. 39–54, 2023.

- [11] M. Lasseck, “Acoustic bird detection with deep convolutional neural networks.” in *DCASE*, 2018, pp. 143–147.
- [12] A. Reiling, W. Mitchell, S. Westberg, E. Balster, and T. Taha, “Cnn optimization with a genetic algorithm,” in *2019 IEEE National Aerospace and Electronics Conference (NAECON)*, 2019, pp. 340–344.